


CUDA Programming

Moreno Marzolla
Dip. di Informatica—Scienza e Ingegneria (DISI)
Università di Bologna

Copyright © 2014, 2017–2026
Moreno Marzolla, Università di Bologna, Italy
<https://www.moreno.marzolla.name/teaching/CP/>



Except where stated otherwise, this work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License (CC BY-SA 4.0). To view a copy of this license, visit <https://creativecommons.org/licenses/by-sa/4.0/> or send a letter to Creative Commons, 543 Howard Street, 5th Floor, San Francisco, California, 94105, USA.

Acknowledgments


- Most of the content of this presentation is from Mark Harris (Nvidia Corporation), “*CUDA C/C++ BASICS*”
 - http://developer.download.nvidia.com/compute/developertrainingmaterials/presentations/cuda_language/Introduction_to_CUDA_C.pptx
- Salvatore Orlando (Univ. Ca' Foscari di Venezia)
- Tim Mattson (Intel Labs)
 - “*Hands-on Intro to CUDA for OpenCL programmers*”
- CUDA C programming guide
 - <http://docs.nvidia.com/cuda/cuda-c-programming-guide/>

CUDA

- CUDA = *Compute Unified Device Architecture*
- Exposes a data-parallel programming model
- CUDA vs OpenCL
 - OpenCL is a portable and open alternative for programming GPUs, CPUs and other architectures
 - OpenCL is based on standard C/C++ plus libraries
 - No vendor lock-in
 - However, OpenCL is being adopted slowly
 - Some implementations not as efficient as CUDA (up to 30% slower)
 - Some implementations are buggy
 - OpenCL programs are much more verbose than CUDA ones
 - Eventually the problems above will be solved

Terminology


- **Host**
 - The **CPU** and its memory (host memory)
- **Device**
 - The **GPU** and its memory (device memory)



Basic concepts


- CUDA separate a program into
 - a **CPU program** (the **host** program), which includes all I/O operations or operation for user interaction, and
 - a **GPU program** (the **device** program), which contains all computations to be executed on the GPU.
- The simplest case of interaction between host program and device program is implemented by
 - a host program that first copies data into the global memory of the GPU
 - the same host program calls device functions to initiate the processing on the GPU

Processing flow



A Programmer's view of a GPU



- A **THREAD** is the minimal unit of work.
 - All the threads execute the same KERNEL function.
 - Threads are grouped in **WARPs** of 32 for scheduling.
- A **BLOCK** is an independent piece of work that can be assigned to a SM.
 - A 1/2/3D array of threads.
 - Max no. of threads in a block: 1024 threads.
- A **GRID** is a piece of work that can be executed by the GPU.
 - A 1/2/3D array of BLOCKS.
- Devices with compute capability ≥ 9.0 allow an intermediate grouping of blocks called *clusters*.



*Streaming
Multiprocessor*

Anatomy of a GPU


*Streaming
Processor*




ming

Threads/Block execution

- There can be many more thread blocks than SMs.
- Runnable blocks are assigned to a SM.
- Scheduling of blocks is non-deterministic.
 - Threads within a block cannot rely on results from other blocks, as they might not be available yet.
 - Threads should not try to synchronize with threads in other blocks.



Blocks and Grid




CUDA thread scheduling

- A **CUDA warp** is a group of 32 CUDA threads that execute simultaneously.
 - Identifiable uniquely by dividing the thread Index by 32.
 - The hardware is most efficiently utilized when all threads in a warp execute instructions from the same program address.
 - If threads in a warp diverge, then some execution pipelines are unused.
 - If threads in a warp access **aligned**, **contiguous** blocks of DRAM, the accesses are coalesced into a single high-bandwidth access.
- A CUDA warp represents the **minimum granularity of efficient SIMT execution**.
 - SIMT = Single Instruction Multiple Threads

Single Instruction Multiple Data

- Individual threads of a warp start together at the same program address.
- Each thread has its own program counter and register state.
 - Each thread is free to branch and execute independently.
 - Provides the MIMD abstraction branch behavior.
- Divergent branches will be executed serially.
 - Threads not following the current branch will be disabled.

```
int x;  
AAA  
if (x > 0) {  
    xxx  
} else {  
    yyy  
}  
BBB
```



Hands-on introduction to CUDA programming

Hello World!

- Standard C that runs on the host.
- The NVIDIA compiler (**nvcc**) can be used to compile programs, even with no device code.

```
/* cuda-hello0.cu */  
#include <stdio.h>  
int main(void)  
{  
    printf("Hello World!\n");  
    return 0;  
}
```

```
$ nvcc cuda-hello0.cu  
$ ./a.out  
Hello World!
```

Running on the HPC cluster

- Prepare a script, e.g., **run-cuda-program.sh**

```
#!/bin/bash

#SBATCH --time=00:01:00
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --gres=gpu:1
#SBATCH --output slurm-%j.out
#SBATCH --partition=rtx2080

echo "== Compiling =="
nvcc cuda-program.cu -o cuda-program

echo "== Running CUDA program =="
./cuda-program

echo "== End of Job =="
```

- **sbatch run-cuda-program.sh**

Hello World! with Device Code

Two new
syntactic
elements

```
/* cuda-hello1.cu */  
#include <stdio.h>  
__global__ void mykernel(void) { }  
  
int main(void)  
{  
    mykernel<<<1,1>>>( );  
    printf("Hello World!\n");  
    return 0;  
}
```

Hello World! with Device Code

```
__global__ void mykernel(void) { }
```

- CUDA/C keyword **__global__** indicates a function that:
 - Runs on the device.
 - Is called from host code.
 - **__global__** functions **must** return **void**.
- **nvcc** separates source code into **host** and **device** components.
 - Device functions (e.g., **mykernel()**) are processed by the NVIDIA compiler.
 - Host functions (e.g., **main()**) are processed by the host compiler.

Hello World! with Device Code

```
mykernel<<<1,1>>>( );
```

- Triple angle brackets mark a call from **host** to **device** code.
 - Also called a “kernel launch”.
 - We’ll return to the parameters (1,1) in a moment.
- That’s all that is required to execute a function on the GPU!

Hello World! with Device Code

```
#include <stdio.h>

__global__ void mykernel(void)
{ }


int main(void) {
    mykernel<<<1,1>>>( );
    printf("Hello World!\n");
    return 0;
}
```

```
$ nvcc cuda-hello1.cu
$ ./a.out
Hello World!
```

- **mykernel ()** does nothing.

Parallel Programming in CUDA/C

- But wait... GPU computing is about massive parallelism!
- We need a more interesting example...
- We'll start by adding two integers and build up to vector addition.



Addition on the Device

- A simple kernel to add two integers.

```
__global__ void add(int *a, int *b, int *c)
{
    *c = *a + *b;
}
```

- As before `__global__` is a CUDA C/C++ keyword meaning:
 - `add()` will execute on the device
 - `add()` will be called from the host

Addition on the Device

- Note the use of pointers for the variables.

```
__global__ void add(int *a, int *b, int *c)
{
    *c = *a + *b;
}
```

- **add()** runs on the **device**, so *a*, *b* and *c* must point to **device** memory.
- We need to allocate memory on the GPU.

Memory Management

- Host and device memory are separate entities:
 - **Device** pointers point to GPU memory.
 - May be passed to/from host code.
 - May **not** be dereferenced in host code.
 - **Host** pointers point to CPU memory.
 - May be passed to/from device code.
 - May **not** be dereferenced in device code.
- Simple CUDA API for handling device memory
 - `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`
 - Similar to the C equivalents `malloc()`, `free()`, `memcpy()`



Addition on the Device: `main()`

```
/* cuda-vecadd0.cu */
int main(void) {
    int a, b, c;           /* host copies of a, b, c */
    int *d_a, *d_b, *d_c; /* device copies of a, b, c */
    const size_t size = sizeof(int);
    /* Allocate space for device copies of a, b, c */
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);
    /* Setup input values */
    a = 2; b = 7;
    /* Copy inputs to device */
    cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);
    /* Launch add() kernel on GPU */
    add<<<1,1>>>(d_a, d_b, d_c);
    /* Copy result back to host */
    cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
    /* Cleanup */
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
    return 0;
}
```

Coordinating Host & Device


- Kernel launches are **asynchronous**.
 - Control returns to the CPU immediately.
- CPU needs to **synchronize** before consuming the results.

cudaMemcpy ()	Blocks the CPU until the copy is complete Copy begins when all preceding CUDA calls have completed
cudaMemcpyAsync ()	Asynchronous wrt to the CPU, i.e., does not block the CPU
cudaDeviceSynchronize ()	Blocks the CPU until all preceding CUDA calls have completed

Running in parallel

Moving to Parallel

- GPU computing is about massive parallelism.
 - So how do we run code in parallel on the device?



- Instead of executing `add()` once, execute N times in parallel.

Vector Addition on the Device

- With **add()** running in parallel we can do vector addition.
- Terminology: each parallel invocation of **add()** is referred to as a **block**.
 - The set of blocks is referred to as a **grid**.
 - Each invocation can refer to its block index using **blockIdx.x**


```
__global__ void add(int *a, int *b, int *c)
{
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];
}
```

- By using **blockIdx.x** to index into the array, each block handles a different index.

Vector Addition on the Device

- On the device, each block can execute in parallel.

```
__global__ void add(int *a, int *b, int *c)
{
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];
}
```



```

/* cuda-vecadd1.cu */
#define N 1024
int main(void) {
    int *a, *b, *c;          /* host copies of a, b, c */
    int *d_a, *d_b, *d_c;    /* device copies of a, b, c */
    const size_t size = N * sizeof(int);
    /* Alloc space for device copies of a, b, c */
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);
    /* Alloc space for host copies of a,b,c and setup input values */
    a = (int *)malloc(size); vec_init(a, N);
    b = (int *)malloc(size); vec_init(b, N);
    c = (int *)malloc(size);
    /* Copy inputs to device */
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    /* Launch add() kernel on GPU with N blocks */
    add<<<N,1>>>>(d_a, d_b, d_c);
    /* Copy result back to host */
    cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
    /* Cleanup */
    free(a); free(b); free(c);
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
    return 0;
}

```

Review

- Difference between **host** and **device**
 - Host $\leftrightarrow$ CPU, device $\leftrightarrow$ GPU
- Using **__global__** to declare a function as device code
 - Executes on the **device**
 - Called from the **host**
- Passing parameters from host code to a device function
- Basic device memory management
 - **cudaMalloc()**
 - **cudaMemcpy()**
 - **cudaFree()**
- Launching parallel kernels
 - Launch N instances of **add()** with **add<<<N,1>>>(...)** ;
 - Use **blockIdx.x** to access block index

Introducing threads

CUDA Threads

- Terminology: a block can be split into parallel threads
- Let's change `add()` to use parallel **threads** instead of parallel **blocks**


```
__global__ void add(int *a, int *b, int *c)
{
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];
}
```

- We use `threadIdx.x` instead of `blockIdx.x`
- Need to make one change in `main()` ...

CUDA Threads

```
__global__ void add(int *a, int *b, int *c)
{
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];
}
```

`blockIdx.x = 0`



```

/* cuda-vecadd2.cu */
#define N 1024
int main(void) {
    int *a, *b, *c;          /* host copies of a, b, c */
    int *d_a, *d_b, *d_c;    /* device copies of a, b, c */
    const size_t size = N * sizeof(int);
    /* Alloc space for device copies of a, b, c */
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);
    /* Alloc space for host copies of a,b,c and setup input values */
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
    /* Copy inputs to device */
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    /* Launch add() kernel on GPU with N threads */
    add<<<1,N>>>(d_a, d_b, d_c);
    /* Copy result back to host */
    cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
    /* Cleanup */
    free(a); free(b); free(c);
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
    return 0;
}

```


Combining threads and blocks

Combining Blocks and Threads

- We have seen parallel vector addition using:
 - Many **blocks** with one **thread** each
 - One **block** with many **threads**
- Let's adapt vector addition to use both blocks and threads
 - Why? We'll come to that...
- First let's discuss data indexing...


Combining Blocks and Threads

- We must somehow assign to each thread a different array element to operate on



Mapping threads to array elements

- No longer as simple as using `blockIdx.x` or `threadIdx.x` alone
- Consider indexing an array with one element per thread, 8 threads per block




- With M threads per block a unique index for each thread is given by:

```
int index = threadIdx.x + blockIdx.x * M;
```


Mapping threads to array elements

- Which thread will operate on the red element?



Array elements

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
---	---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----

```
int index = threadIdx.x + blockIdx.x * M;  
          =      4      +      2      * 8;  
          = 20;
```

Vector Addition with Blocks and Threads

- Use the built-in variable `blockDim.x` to get the number of threads per block

```
int index = threadIdx.x + blockIdx.x * blockDim.x;
```

- Combined version of `add()` to use parallel threads and parallel blocks:

```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[index] = a[index] + b[index];  
}
```

- What changes need to be made in `main()`?

Addition with blocks and threads

```
#define N (2048*2048)
#define BLKDIM 1024

int main(void) {
    ...
    /* Launch add() kernel on GPU */
    add<<<N/BLKDIM, BLKDIM>>>(d_a, d_b, d_c);
    ...
}
```

*Number of
blocks*

*Number of threads
per block*

- However, the problem size might not be multiple of the block size...

Handling arbitrary vector sizes

- Avoid accessing beyond the end of the array

```
__global__ void add(int *a, int *b, int *c, int n) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    if (index < n) {  
        c[index] = a[index] + b[index];  
    }  
}
```

- Update kernel launch

```
add<<< (N + BLKDIM-1)/BLKDIM, BLKDIM>>> (d_a, d_b, d_c, N);
```

- See [cuda-vecadd3.cu](#)

Review

- Launching parallel kernels
 - Launch $\sim N$ copies of `add()` with
`add<<<(N + BLKDIM-1)/BLKDIM, BLKDIM>>>(...);`
 - Use `blockIdx.x` to access block index
 - Use `threadIdx.x` to access thread index within block
- Assign array elements to threads:

```
int index = threadIdx.x + blockIdx.x * blockDim.x;
```


Why bother with Blocks?

- Blocks seem unnecessary
 - They add a level of complexity
 - What do we gain?
- Threads within the same block have mechanisms to:
 - Share information efficiently
 - Synchronize
- To look closer, we need a new example...

Cooperating threads


1D Stencil

- Consider applying a stencil to a 1D array of elements
 - Each output element is the sum of input elements within a given **radius**
- If RADIUS is 3, then each output element is the sum of 7 input elements
 - The first and last RADIUS elements of the output array are not computed



Implementing Within a Block


- Each thread processes one output element
 - `blockDim.x` elements per block
- Input elements are read several times
 - With radius 3, each input element is read **seven times**
 - With radius R , each input element is read $(2R+1)$ times



Sharing Data Between Threads


- Global memory accesses are likely to cause a **bottleneck** due to the limited memory bandwidth
- Within a block, threads can share data via **shared memory**
 - Extremely fast on-chip memory, user-managed
 - Think of it as a user-managed local cache
- Declare using **__shared__**, allocated per thread-block
- Data is not visible to threads in other blocks

CUDA memory model



Implementing with shared memory


- Cache data in shared memory
 - Read (`blockDim.x + 2 × radius`) input elements from **global** memory to **shared** memory
 - Compute `blockDim.x` output elements
 - Write `blockDim.x` output elements to global memory
 - Each block needs a **ghost area (halo)** of *radius* elements at each boundary



Implementing with shared memory


- Let us make a few simplifying assumptions
 - The array length is a multiple of the thread block size
 - The input (and output) array already includes an halo of **2*RADIUS** elements
 - The halo is ignored for the output array
- Idea
 - Each thread block keeps a local cache of **blockDim.x + 2*RADIUS** elements
 - Each thread copies one element from the global array to the local cache
 - The first **RADIUS** threads also take care of filling the halo

Implementing with shared memory




```
__shared__ int temp[BLKDIM + 2 * RADIUS];
const int gindex = threadIdx.x + blockIdx.x * blockDim.x + RADIUS;
const int lindex = threadIdx.x + RADIUS;
/* ... */
temp[lindex] = in[gindex];
if (threadIdx.x < RADIUS) {
    temp[lindex - RADIUS] = in[gindex - RADIUS];
    temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];
}
```

Implementing with shared memory



```
__shared__ int temp[BLKDIM + 2 * RADIUS];  
const int gindex = threadIdx.x + blockIdx.x * blockDim.x + RADIUS;  
const int lindex = threadIdx.x + RADIUS;  
/* ... */  
temp[lindex] = in[gindex];  
if (threadIdx.x < RADIUS) {  
    temp[lindex - RADIUS] = in[gindex - RADIUS];  
    temp[lindex + BLKDIM] = in[gindex + BLKDIM];  
}
```

Implementing with shared memory



```
__shared__ int temp[BLKDIM + 2 * RADIUS];  
const int gindex = threadIdx.x + blockIdx.x * blockDim.x + RADIUS;  
const int lindex = threadIdx.x + RADIUS;  
/* ... */  
temp[lindex] = in[gindex];  
if (threadIdx.x < RADIUS) {  
    temp[lindex - RADIUS] = in[gindex - RADIUS];  
    temp[lindex + BLKDIM] = in[gindex + BLKDIM];  
}
```



Stencil kernel (does not work!)

Wrong!

```
__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLKDIM + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x + RADIUS;
    int lindex = threadIdx.x + RADIUS;
    int result = 0, offset;
    /* Read input elements into shared memory */
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        temp[lindex + blockDim.x] = in[gindex + blockDim.x];
    }
    /* Apply the stencil */
    for (offset = -RADIUS ; offset <= RADIUS ; offset++) {
        result += temp[lindex + offset];
    }
    /* Store the result */
    out[gindex] = result;
}
```

The problem


- All threads are not necessarily fully synchronized
- Suppose that thread (`blockDim.x - 1`) reads the halo before thread 0 has fetched it
 - Data race!



The solution: `__syncthreads ()`

- Synchronizes all threads within a block.
 - Used to prevent RAW / WAR / WAW hazards.
- All threads must reach the barrier.
 - In conditional code, the condition must be uniform across the block.

Stencil kernel that works

```
__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLKDIM + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x + RADIUS;
    int lindex = threadIdx.x + RADIUS;
    int result = 0, offset;
    /* Read input elements into shared memory */
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        temp[lindex + blockDim.x] = in[gindex + blockDim.x];
    }
    __syncthreads(); 
    /* Apply the stencil */
    for (offset = -RADIUS ; offset <= RADIUS ; offset++) {
        result += temp[lindex + offset];
    }
    /* Store the result */
    out[gindex] = result;
}
```

See [cuda-stencil1d-shared.c](#)

Review

- Use **__shared__** to declare a variable/array in shared memory
 - Data is shared between threads in a block
 - Not visible to threads in other blocks
- Use **__syncthreads()** as a barrier
 - Use to prevent data hazards

Managing the device

Timing CUDA kernels

- You can use the timing routines provided in `hpc.h`
- However, since kernel invocations are **asynchronous**, you must call `cudaDeviceSynchronize()` to wait the kernel to complete execution

```
#include "hpc.h"
...
double tstart, tend;

tstart = hpc_gettime();
mykernel<<<X, Y>>>( ); /* kernel invocation */
cudaDeviceSynchronize(); /* wait for kernel to finish */
tend = hpc_gettime();

printf("Elapsed time %f\n", tend - tstart);
```

The `__device__` qualifier

- The `__device__` qualifier defines functions that
 - execute on the device
 - can be called from device code only
- `__device__` functions are inlined, so they can return a value

```
__device__ float cuda_fmaxf(float a, float b)
{
    return (a>b ? a : b);
}

__global__ void my_kernel( float *v, int n )
{
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i<n) {
        v[i] = cuda_fmaxf(1.0, v[i]);
    }
}
```


The `__host__` qualifier

- (Optional) denotes a function that
 - is executed on the host
 - can be called from host code only
 - default behavior when neither `__global__` nor `__device__` are specified
- You can use both the `__host__` and `__device__` qualifiers for the same function
 - The compiler produces two versions of the function: one for the GPU and one for the CPU

```
__host__ __device__ float my_fmaxf(float a, float b)
{
    return (a>b ? a : b);
}
```

Recap function declaration

	Executed on:	Only callable from:
<code>__device__ float deviceFunc()</code>	Device	Device
<code>__host__ float hostFunc()</code>	Host	Host
<code>__global__ void kernelFunc()</code>	Device	Host


The `__device__` qualifier

- The `__device__` qualifier can be applied to variables
 - easy way to statically allocate variables on the device
 - can be applied to **global variables only**
 - `cudaMemcpyToSymbol()` / `cudaMemcpyFromSymbol()` must be used to copy data from/to `__device__` variables
 - will see an example in the lab

```
#define BLKDIM 1024

__device__ float d_buf[BLKDIM];

int main( void )
{
    float buf[BLKDIM];
    ...
    cudaMemcpyToSymbol(d_buf, buf, BLKDIM*sizeof(float));
    ...
    cudaMemcpyFromSymbol(buf, d_buf, BLKDIM*sizeof(float));
}
```



Reporting Errors

- All CUDA API calls return an error code (**cudaError_t**)
 - Error in the API call itself, OR
 - Error in an earlier asynchronous operation (e.g., kernel)
 - **cudaSuccess** means no error
- Get the error code for the last error:
`cudaError_t cudaGetLastError(void)`
- Get a string to describe the error:
`char *cudaGetErrorString(cudaError_t)`
`printf("%s\n", cudaGetErrorString(cudaGetLastError()));`

Reporting Errors

- Some useful macros defined in **hpc.h**
 - **cudaSafeCall(Exp)** execute Exp and checks for return status
 - Exp can be any CUDA function returning an error code, e.g., `cudaMalloc`, `cudaMemcpy`, ...
 - **cudaCheckError()** checks error code from last CUDA operation
 - Usually, a kernel call
 - **cudaCheckError()** calls **cudaDeviceSynchronize()**




```
cudaSafeCall( cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice) );  
my_kernel<<< 1, 1 >>>(d_a); cudaCheckError();  
cudaSafeCall( cudaMemcpy(h_a, d_a, size, cudaMemcpyDeviceToHost) );
```

Higher Block Dimensions

Matrix-Matrix multiply


- Multiply two $n \times n$ matrices p and q

```
void matmul( const float *p, const float* q,
             float *r, int n)
{   int i, j, k; float v;
    for (i=0; i<n; i++) {
        for (j=0; j<n; j++) {
            v = 0.0;
            for (k=0; k<n; k++) {
                v += p[i*n + k] * q[k*n + j];
            }
            r[i*n + j] = v;
        }
    }
}
```



Matrix-Matrix multiply with thread blocks

- Decompose the result matrix r into square blocks
- Assign each block to a thread block
 - All threads of the thread block compute one element of the result matrix



Setting up the thread blocks

- The `dim3` data type is used to define a one-, two-, or three-dimensional “size” for a block or grid
 - `dim3 blk(3)`
 - defines a variable “blk” representing a 3x1x1 block
 - `dim3 blk(3, 4)`
 - defines a variable “blk” representing a 3x4x1 block
 - `dim3 blk(3, 4, 7)`
 - defines a variable “blk” representing a 3x4x7 block

Examples

- Launch 3 blocks, 16 threads per block (1D)

```
mykernel<<<3, 16>>>( );
```

- The same

```
dim3 grid(3);  
dim3 block(16);  
mykernel<<<grid, block>>>( );
```

- Launch (16×4) blocks, $(8 \times 8 \times 8)$ threads per block

```
dim3 grid(16, 4);    /* 2D */  
dim3 block(8, 8, 8); /* 3D */  
mykernel<<<grid, block>>>( );
```

Kernel invocation

- Setup and kernel invocation


```
/* cuda-matmul.cu */
#define BLKDIM 32

int main( void )
{
    ...
    dim3 block(BLKDIM, BLKDIM);
    dim3 grid((N+BLKDIM-1)/BLKDIM, (N+BLKDIM-1)/BLKDIM);
    ...
    /* Launch matmul() kernel on the GPU */
    matmul<<<grid, block>>>(d_p, d_q, d_r, N);
    ...
}
```

The matmul kernel

- Each thread computes one element $r[i][j]$ of matrix r

```
const int i = blockIdx.y * blockDim.y + threadIdx.y;  
const int j = blockIdx.x * blockDim.x + threadIdx.x;
```



The matmul kernel


- The kernel function:

```
__global__ void matmul( const float *p, const float *q,
                        float *r, int n )
{
    const int i = blockIdx.y * blockDim.y + threadIdx.y;
    const int j = blockIdx.x * blockDim.x + threadIdx.x;
    int k;
    float val = 0.0;
    if ( i < n && j < n ) {
        for (k=0; k<n; k++) {
            val += p[i*n + k] * q[k*n + j];
        }
        r[i*n + j] = val;
    }
}
```

Matrix-Matrix multiply



- Multiply two $n \times n$ matrices p and q

```
void matmul( const float *p, const float* q,
             float *r, int n)
{   int i, j, k; float v;
    for (i=0; i<n; i++) {
        for (j=0; j<n; j++) {
            v = 0.0;
            for (k=0; k<n; k++) {
                v += p[i*n + k] * q[k*n + j];
            }
            r[i*n + j] = v;
        }
    }
}
```





How many times are the elements of matrix p accessed?

How many times are the elements of matrix q accessed?




Reducing the memory pressure

- To reduce the number of reads from global memory we use shared memory to cache the data needed by each block to compute its portion of the result
- This requires $(2 \times BLKDIM \times n)$ elements, which might exceed the amount of shared memory allowed by the device






CUDA program



Reducing the memory pressure





- The solution: divide the $(BLKDIM \times n)$ stripes of p and q into square blocks of $(BLKDIM \times BLKDIM)$ elements each
- Operate on two blocks (for p and q) at a time

$$R = P_1 \times Q_1 + P_2 \times Q_2 + P_3 \times Q_3$$




How the new kernel works

- For each m :




How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$







How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel




How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$




How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$







How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$




How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$




How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$



How the new kernel works

- For each m :
 - Copy blocks from p and q into shared memory $local_p$ and $local_q$ in parallel
 - compute the matrix product $local_p \times local_q$ in parallel
 - $m \leftarrow m + BLKDIM$






The new matmul kernel

```
__global__ void matmulb( const float *p, const float *q,  
                        float *r, int n )  
{  
    __shared__ float local_p[BLKDIM][BLKDIM];  
    __shared__ float local_q[BLKDIM][BLKDIM];  
    const int bx = blockIdx.x; const int by = blockIdx.y;  
    const int tx = threadIdx.x; const int ty = threadIdx.y;  
    const int i = by * BLKDIM + ty;  
    const int j = bx * BLKDIM + tx;  
    float v = 0.0;  
    for (int m = 0; m < n; m += BLKDIM) {  
        local_p[ty][tx] = p[i*n + (m + tx)];  
        local_q[ty][tx] = q[(m + ty)*n + j];  
        __syncthreads();  
        for (int k = 0; k < BLKDIM; k++) {  
            v += local_p[ty][k] * local_q[k][tx];  
        }  
        __syncthreads();  
    }  
    r[i*n + j] = v;  
}
```

```

local_p[ty][tx] = p[i*n + (m + tx)];
local_q[ty][tx] = q[(m + ty)*n + j];


```



The new matmul kernel

- The setup and kernel invocation remain the same
- See [cuda-matmul.cu](https://nvidia.github.io/cuda-matmul)

How are threads grouped into warps?



```
__global__ myKernel( void )  
{  
    const int WARP_SIZE = 32;  
    const int tid = threadIdx.x + threadIdx.y * blockDim.x;  
    const int warp_id = tid / WARP_SIZE;  
    // threads with the same warp_id belong to the same warp  
    ...  
}
```


Reduction on the GPU

Problem statement

- Given a nonempty array of floats or ints, compute the sum of all elements of the array
- Basic idea
 - Decompose the problem across thread blocks
 - Each thread block computes a partial reduction
 - The CPU completes the reduction

Solution #0 (naive)

- Thread 0 of each block computes the local sum



Solution #0 kernel

```
#define BLKDIM 1024
#define N_OF_BLOCKS 1024
#define N ((N_OF_BLOCKS) * (BLKDIM))

__device__ int d_sums[N_OF_BLOCKS];
int h_sums[N_OF_BLOCKS];

__global__ void sum( int *a, int n )
{
    __shared__ int temp[BLKDIM];
    int lindex = threadIdx.x;
    int bindex = blockIdx.x;
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    temp[lindex] = a[gindex];
    __syncthreads();
    if ( 0 == lindex ) {
        int my_sum = 0;
        for (int i=0; i<blockDim.x; i++) {
            my_sum += temp[i];
        }
        d_sums[bindex] = my_sum;
    }
}
```


This kernel requires that n is a multiple of BLKDIM.

Shared memory is not useful here; it will be useful in the other versions of this kernel

See [cuda-reduction0.cu](#)

Solution #1 (better)

- All threads within each block cooperate to compute the local sum



Solution #1 kernel

```
__global__ void sum( int *a, int n )
{
    __shared__ int temp[BLKDIM];
    const int lindex = threadIdx.x;
    const int bindex = blockIdx.x;
    const int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int bsize = blockDim.x / 2;
    temp[lindex] = a[gindex];
    __syncthreads();
    while ( bsize > 0 ) {
        if ( lindex < bsize ) {
            temp[lindex] += temp[lindex + bsize];
        }
        bsize = bsize / 2;
        __syncthreads();
    }
    if ( 0 == lindex ) {
        d_sums[bindex] = temp[0];
    }
}
```


This kernel requires that

- BLKDIM is a power of two.
- n is a multiple of BLKDIM.

See [cuda-reduction1.cu](#)

Solution #2 (one CUDA block)

- All threads within each block cooperate to compute the local sum



Solution #2 kernel


```
__global__ void sum( int *a, int n, int *d_result )
{
    __shared__ int temp[BLKDIM];
    const int lindex = threadIdx.x;
    const int bindex = blockIdx.x;
    const int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int bsize = blockDim.x / 2;
    temp[lindex] = 0;
    for (int i = lindex; i < n; i += blockDim.x)
        temp[lindex] += a[i];
    syncthreads();
    while ( bsize > 0 ) {
        if ( lindex < bsize ) {
            temp[lindex] += temp[lindex + bsize];
        }
        bsize = bsize / 2;
        syncthreads();
    }
    if ( 0 == lindex ) {
        *d_result = temp[0];
    }
}
```

This kernel requires that
BLKDIM is a power of two.

See cuda-reduction4.cu

Memory Access Optimization Techniques for GPUs

CUDA memory model




Memory access patterns

- Global memory can be accessed through 32, 64 or 128-byte transactions.
 - So, if a thread just needs a single float (4B), this results in 32B, 64B or 128B being moved.
 - Exact details vary depending on the CUDA capability.
- Memory transfers must be naturally aligned (at 32, 64 or 128-byte boundary).
 - However, the hardware allows memory accesses at any address.
- The GPU can pack together (coalesce) memory accesses of **threads of the same warp** (32 threads)
 - Examples follow.


Caching load

- Warp requests 32 aligned, consecutive 4-byte words
- Addresses fall within 1 cache-line
 - Warp needs 128 bytes
 - 128 bytes move across the bus on a miss
 - Bus utilization: 100%
 - Transactions: 1




Caching load

- Warp requests 32 aligned, permuted 4-byte words
- Addresses fall within 1 cache-line
 - Warp needs 128 bytes
 - 128 bytes move across the bus on a miss
 - Bus utilization: 100%
 - Transactions: 1




Caching load

- Warp requests 32 misaligned, consecutive 4-byte words
- Addresses fall within 2 cache-lines
 - Warp needs 128 bytes
 - 256 bytes move across the bus on misses
 - Bus utilization: 50%
 - Transactions: 2




Caching load

- All threads in a warp request the same 4-byte word
- Addresses fall within a single cache-line
 - Warp needs 4 bytes
 - 128 bytes move across the bus on a miss
 - Bus utilization: 3.125%




Caching load

- Warp requests 32 scattered 4-byte words
- Addresses fall within N cache-lines
 - Warp needs 128 bytes
 - $N \times 128$ bytes move across the bus on a miss
 - Bus utilization: $128 / (N \times 128)$



Example


- See [cuda-image-rotation.cu](#)




Example

- See cuda-image-rotation.cu

Device memory




Shared memory



Example


- See cuda-image-rotation.cu



Example

- Array of Structures (AoS) containing three floats.
 - Think x, y, z as coordinates of a set of points.
- Each memory transaction read data that is not used.
 - Low efficiency.


```
typedef struct {  
    float x;  
    float y;  
    float z;  
} point3d;  
  
__global__ void  
compute( point3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points[i].x;  
    float y = points[i].y;  
    float z = points[i].z;  
    /* use x, y, z ... */  
}
```



Example

- Array of Structures (AoS) containing three floats.
 - Think x, y, z as coordinates of a set of points.
- Each memory transaction read data that is not used.
 - Low efficiency.


```
typedef struct {  
    float x;  
    float y;  
    float z;  
} point3d;  
  
__global__ void  
compute( point3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points[i].x;  
    float y = points[i].y;  
    float z = points[i].z;  
    /* use x, y, z ... */  
}
```



Example

- Array of Structures (AoS) containing three floats.
 - Think x, y, z as coordinates of a set of points.
- Each memory transaction read data that is not used.
 - Low efficiency.


```
typedef struct {  
    float x;  
    float y;  
    float z;  
} point3d;  
  
__global__ void  
compute( point3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points[i].x;  
    float y = points[i].y;  
    float z = points[i].z;  
    /* use x, y, z ... */  
}
```



Example

- Structure of Arrays (SoA) containing three arrays of floats.
- Memory accesses can now be coalesced.


```
typedef struct {  
    float *x;  
    float *y;  
    float *z;  
} points3d;  
  
__global__ void  
compute( points3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points->x[i];  
    float y = points->y[i];  
    float z = points->z[i];  
    /* use x, y, z ... */  
}
```



Example

- Structure of Arrays (SoA) containing three arrays of floats.
- Memory accesses can now be coalesced.

```
typedef struct {  
    float *x;  
    float *y;  
    float *z;  
} points3d;  
  
__global__ void  
compute( points3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points->x[i];  
    float y = points->y[i];  
    float z = points->z[i];  
    /* use x, y, z ... */  
}
```




Example

- Structure of Arrays (SoA) containing three arrays of floats.
- Memory accesses can now be coalesced.

```
typedef struct {  
    float *x;  
    float *y;  
    float *z;  
} points3d;  
  
__global__ void  
compute( points3d *points, int n )  
{  
    int i = threadIdx.x;  
    float x = points->x[i];  
    float y = points->y[i];  
    float z = points->z[i];  
    /* use x, y, z ... */  
}
```

Addresses from a warp




GPU memory optimization guidelines

- Try to coalesce memory accesses.
 - Align starting address (may require padding).
 - A warp should access contiguous memory locations.
 - Restructure your data into a Structure of Arrays (SoA).
- Have enough concurrent accesses to saturate the bus.
 - Launch enough threads to maximize throughput.

Performance Evaluation of CUDA programs

Performance Evaluation

- The usual concept of speedup can not be used.
 - The program has no control over the number of CUDA cores used.
 - The hardware multiplexes CUDA threads to CUDA cores.
- We need different metrics, e.g.
 - **Throughput**: number of operations/seconds as a function of the input size.
 - **Speedup vs efficient (i.e., parallel) CPU implementation.**



Source: <https://moderngpu.github.io/scan.html>

Performance Evaluation

- **Throughput**
 - Number of operations/seconds as a function of input size.
 - See next slide.
- **Speedup vs efficient CPU implementation**
 - Computed as the ratio:

$$\text{CUDA Speedup} = \frac{\text{OpenMP wall-clock time}}{\text{CUDA wall-clock time}}$$

using the maximum number of CPU cores available.

- Comparing a CUDA program with a serial CPU implementation is wrong.

Estimating Throughput

$$\text{Throughput} = \frac{\text{N. of operations}}{\text{wall-clock time}}$$

- **Number of operations**
 - Estimated with the (asymptotic) work performed by the algorithm
 - E.g.,
 - $O(n^2)$ algorithm.
 - Input size $n = 1000$.
 - $\rightarrow$ N. of operations $= 1000^2 = 10^6$.
- **Wall-clock time**
 - Measured from the running program.
- Throughput is measured in **[M,G]ops/sec**.
- See example.